

PCA, SVD and Mahalanobis distance

Christopher Gillies

11/16/2017

Multivariate normal distribution

$$f(x) = \frac{1}{\sqrt{(2\pi)^n |\Sigma|}} \exp \left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right) \quad (1)$$

If we assume the data are zero-centered then:

$$f(x) = \frac{1}{\sqrt{(2\pi)^n |\Sigma|}} \exp \left(-\frac{1}{2} x^T \Sigma^{-1} x \right) \quad (2)$$

Generate a random sample

```
require(MASS)
```

```
## Loading required package: MASS
```

```
require(ggplot2)
```

```
## Loading required package: ggplot2
```


```
## Warning: package 'ggplot2' was built under R version 3.3.2
```

```
S = matrix(c(1,0.75,0.75,2),ncol=2)
```

```
x = mvrnorm(n = 1000, mu=c(0,0), Sigma=S)
```

```
ggplot(data.frame()) + geom_point(aes(x=x[,1],y=x[,2])) + scale_x_continuous(limits=c(-4,4)) + scale_y_
```

```
## Warning: Removed 4 rows containing missing values (geom_point).
```



Perform PCA

```
S.sample = cov(x)


e.decomp = eigen(S.sample)

e.decomp$vectors %*% diag(e.decomp$values) %*% t(e.decomp$vectors)

##           [,1]      [,2]
## [1,] 1.0141466 0.7407353
## [2,] 0.7407353 2.0021503

x.proj = x %*% e.decomp$vectors
ggplot(data.frame()) + geom_point(aes(x=x.proj[,1],y=x.proj[,2])) + scale_x_continuous(limits=c(-4,4))

## Warning: Removed 8 rows containing missing values (geom_point).
```



```
e.decomp$values
```

```
## [1] 2.3985004 0.6177964
```

```
var(x.proj[,1])
```

```
## [1] 2.3985
```

```
var(x.proj[,2])
```

```
## [1] 0.6177964
```

Notice that the variance of the projected data matches the eigenvalues of the covariance matrix. The projection of x onto its principal components simply rotates the data.

```
svd.sample.cov = svd(S.sample)
```


```
x.center = scale(x,scale = FALSE)
```

```
svd.x = svd(x.center)
```


```
x.proj.2 = x.center %*% svd.x$v
```

```
x.proj.3 = x.center %*% svd.sample.cov$v
```

```
plot(x.proj.2[,1],x.proj.3[,1])
```



```
plot(x.proj.2[,2],x.proj.3[,2])
```



```
svd.x$d^2 / (1000 - 1)
```

```
## [1] 2.3985004 0.6177964
```

```
svd.sample.cov$d
```

```
## [1] 2.3985004 0.6177964
```

The same eigenvectors and eigenvalues are computed from the matrix $x.\text{center}$ and the covariance matrix. The singular values σ_i of x and the eigenvalues λ_i of $\text{cov}(x)$ are related as follows:

$$\lambda_i = \frac{\sigma_i^2}{n-1} \quad (3)$$

$$\text{COV}[X] = \frac{1}{n-1} X^T X \quad (4)$$

$$X = U \Sigma V^T \quad (5)$$

$$X^T X = (U \Sigma V^T)^T (U \Sigma V^T) \quad (6)$$

$$X^T X = V \Sigma U^T U \Sigma V^T \quad (7)$$

$$X^T X = V \Sigma \Sigma V^T = V \Sigma^2 V^T \quad (8)$$

$$\frac{1}{n-1} X^T X = \frac{1}{n-1} V \Sigma^2 V^T \rightarrow \frac{1}{n-1} \Sigma^2 = \Lambda \quad (9)$$

where Λ is the diagonal matrix of eigenvalues of $\text{COV}[X]$.

This is the same formula as above for the relationship between the singular values of X and the eigenvalues of its covariance matrix.

What happens if we scale X before running PCA?

```
x.scaled = scale(x)
```

```
cov.scaled.x = cov(x.scaled)
cov.scaled.x
```

```
##           [,1]      [,2]
## [1,] 1.0000000 0.5198336
## [2,] 0.5198336 1.0000000
```


```
S.sample
```

```
##           [,1]      [,2]
## [1,] 1.0141466 0.7407353
## [2,] 0.7407353 2.0021503
```

```
ggplot(data.frame()) + geom_point(aes(x=x[,1],y=x[,2])) + scale_x_continuous(limits=c(-4,4)) + scale_y_
```


```
## Warning: Removed 4 rows containing missing values (geom_point).
```

Before scaling



```
ggplot(data.frame()) + geom_point(aes(x=x.scaled[,1],y=x.scaled[,2])) + scale_x_continuous(limits=c(-4,
```


After scaling



```
scaled.svd = svd(cov.scaled.x)

x.scaled.proj = x.scaled %*% scaled.svd$v
ggplot(data.frame()) + geom_point(aes(x=x.scaled.proj[,1],y=x.scaled.proj[,2])) + scale_x_continuous(lim=
```

After projection scaled x



```
svd.sample.cov$d
```

```
## [1] 2.3985004 0.6177964
```

```
scaled.svd$d
```

```
## [1] 1.5198336 0.4801664
```

Data reconstruction

```
x.projection = svd.sample.cov$v %*% t(x.center)
x.reconstruction = t(svd.sample.cov$v %*% x.projection)
x.reconstruction.partial = t(svd.sample.cov$v %*% rbind(x.projection[1,], 0))
```

```
cor(x.reconstruction[,1],x.center[,1])
```

```
## [1] 1
```

```
cor(x.reconstruction[,2],x.center[,2])
```

```
## [1] 1
```

```
cor(x.reconstruction[,1],x.reconstruction.partial[,1])
```

```
## [1] 0.725543
```

```
cor(x.reconstruction[,2],x.reconstruction.partial[,2])
```

```
## [1] 0.9650488
```


Notice that the reconstruction works perfectly, and the partial reconstruction works fairly well.

Whitening transform

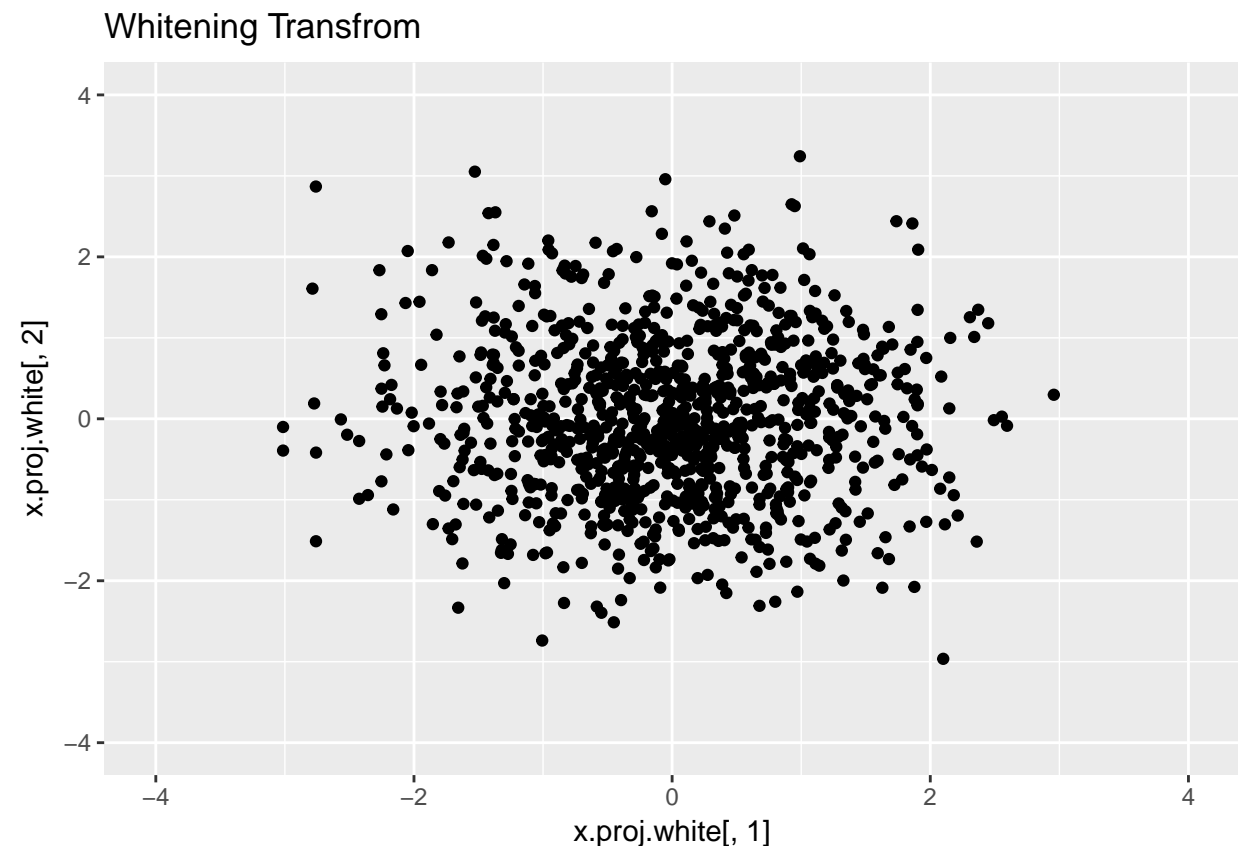
```
x.proj.white = x.center %*% svd.sample.cov$v %*% solve(diag(sqrt(svd.sample.cov$d)))  
var(x.proj.white[,1])
```

```
## [1] 1
```

```
var(x.proj.white[,2])
```


```
## [1] 1
```

```
ggplot(data.frame()) + geom_point(aes(x=x.proj.white[,1],y=x.proj.white[,2])) + scale_x_continuous(limi
```



```
w_sqrt = svd.sample.cov$v %*% diag(sqrt(svd.sample.cov$d)) %*% t(svd.sample.cov$v)  
w_inv_sqrt = solve(w_sqrt)  
w_inv_sqrt.2 = t(svd.sample.cov$v) %*% solve(diag(sqrt(svd.sample.cov$d))) %*% svd.sample.cov$v  
x.proj.white.2= x.center %*% w_inv_sqrt.2  
ggplot(data.frame()) + geom_point(aes(x=x.proj.white.2[,1],y=x.proj.white.2[,2])) + scale_x_continuous(limi
```

Whitening Transform 2



```
cov(x.proj.white.2)
```

```
##           [,1]           [,2]
## [1,]  1.000000e+00 -3.038469e-16
## [2,] -3.038469e-16  1.000000e+00
```

```
cov(x.proj.white)
```

```
##           [,1]           [,2]
## [1,]  1.000000e+00 -8.979573e-17
## [2,] -8.979573e-17  1.000000e+00
```

The above are two different whitening transforms. Please note that:

$$\text{COV}[X] = V\Lambda V^T \quad (10)$$

In a diagonal matrix Λ we can take the square roots by just taking the square root of the diagonal elements

$$\text{COV}[X] = (V\Lambda^{1/2}V^T)^2 = V\Lambda^{1/2}(V^TV)\Lambda^{1/2}V^T = V\Lambda V^T \quad (11)$$

Since, $V^TV = I$. That is $V^T = V^{-1}$.

$$\text{COV}[X]^{1/2} = V\Lambda^{1/2}V^T \quad (12)$$

Now $(ABC)^{-1} = C^{-1}B^{-1}A^{-1}$. Therefore

$$\text{COV}[X]^{-1/2} = (V\Lambda^{1/2}V^T)^{-1} = V^{-1}\Lambda^{-1/2}(V^{-1})^{-1} = V^T\Lambda^{-1/2}V \quad (13)$$

Since $V^T = V^{-1}$.

Mahalanobis distance²

The Mahalanobis of a centered vector from the origin is:

$$D_M(x)^2 = x^T S^{-1} x \quad (14)$$

where S is the covariance matrix of the vector x .

```
t(x.center[1,]) %%% solve(S.sample) %%% x.center[1,]

##           [,1]
## [1,] 0.3627102

sum(x.proj.white[1,]^2)

## [1] 0.3627102

sum(x.proj.white.2[1,]^2)

## [1] 0.3627102

x.1.proj.scaled = x.center[1,] %%% svd.sample.cov$v %%% solve(diag(sqrt(svd.sample.cov$d)))
sum(x.1.proj.scaled^2)

## [1] 0.3627102
```

Mahalanobis distance² is the equal to the $\|x\|^2$ in the Whitened space.

The norm of a vector

$$\|x\|^2 = \sum x_i^2 = x^T x \quad (15)$$

$$(x^T)^T = x \quad (16)$$

$$D_M(x)^2 = x^T \text{COV}[X]^{-1} x \quad (17)$$

$$\text{COV}[X]^{-1} = (V \Lambda V^T)^{-1} = V^T \Lambda^{-1} V \quad (18)$$

The Mahalanobis distance² can be written as

$$\begin{aligned} D_M(x)^2 &= x^T V^T \Lambda^{-1} V x = x^T V^T \Lambda^{-1/2} \Lambda^{-1/2} V x = (x^T V^T \Lambda^{-1/2}) (\Lambda^{-1/2} V x) = (\Lambda^{-1/2} V x)^T (\Lambda^{-1/2} V x) \\ &= (\Lambda^{-1/2} V x)^T (\Lambda^{-1/2} V x) = \|\Lambda^{-1/2} V x\|^2 = \|x^T V^T \Lambda^{-1/2}\|^2 = D_M(x)^2 \end{aligned} \quad (19)$$

The final equality $\|x^T V^T \Lambda^{-1/2}\|^2$ shows that the $D_M(x)^2$ is the same as projecting the x onto its principal components, then scaling each axis by the square root of its eigenvalue (if the eigenvalue is the variance then the $\sqrt{\text{eigenvalue}}$ is like the standard deviation) and finally taking the norm of the scaled projected x . A eigenvalue is the variance of the data projected onto its corresponding eigenvector, so to scale it you divide by the standard deviation.

What about PCA and SVD in the case where there are more variables than observations?

Generate random samples. Assume we have $n = 100$ subjects and $m = 200$ snps.

```
n = 100
m = 50000
z = rbinom(n,size = 2, prob = 0.5)
X = matrix(rep(0,n * m),ncol=n)
snp_afs_0 = runif(m,0.01,0.5)
snp_afs_1 = runif(m,0.25,0.75)
snp_afs_2 = c(snp_afs_0[1:(m/2)],snp_afs_1[(m/2 + 1):m])
for(i in 1:n) {
  x_i = NULL
  if(z[i] == 0) {
    x_i = rbinom(m,size=2,prob=snp_afs_0)
  } else if(z[i] == 1) {
    x_i = rbinom(m,size=2,prob=snp_afs_1)
  } else {
    x_i = rbinom(m,size=2,prob=snp_afs_2)
  }
  X[,i] = x_i
}
```

Standardize by subtracting off the row means and dividing by the standard deviation. Also compute covariance matrix.

```
X_std = t(scale(t(X)))
ind_cov = cov(X_std)
ind_cor_x = cor(X)

#approx covariance matrix
Xt_X = t(X_std) %*% X_std * 1/(n-1)

dim(ind_cov)

## [1] 100 100

dim(Xt_X)

## [1] 100 100

sign(Xt_X[1:10,1:10]) == sign(ind_cov[1:10,1:10])
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
## [1,] TRUE FALSE TRUE TRUE TRUE FALSE TRUE TRUE FALSE TRUE
## [2,] FALSE TRUE FALSE TRUE TRUE TRUE FALSE TRUE TRUE TRUE
## [3,] TRUE FALSE TRUE TRUE TRUE FALSE TRUE TRUE FALSE TRUE
## [4,] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [5,] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [6,] FALSE TRUE FALSE TRUE TRUE TRUE FALSE TRUE TRUE TRUE
## [7,] TRUE FALSE TRUE TRUE TRUE FALSE TRUE TRUE FALSE TRUE
## [8,] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [9,] FALSE TRUE FALSE TRUE TRUE TRUE FALSE TRUE TRUE TRUE
## [10,] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

```
Xt_X[1:10,1:10] / ind_cov[1:10,1:10]
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,]  511.1433 -2299.8931  552.5964  741.8059  718.7674 -3109.7082
## [2,] -2299.8931  625.4373 -2473.9409 1348.9850 1342.6063  1309.7657
## [3,]  552.5964 -2473.9409  511.1046  735.6752  728.6359 -2101.6812
## [4,]  741.8059 1348.9850  735.6752  533.5258 1334.3593  1408.8509
## [5,]  718.7674 1342.6063  728.6359 1334.3593  531.9334  1354.4483
## [6,] -3109.7082 1309.7657 -2101.6812 1408.8509 1354.4483  629.4193
## [7,]  554.9157 -2286.1969  554.8381  738.0696  711.8100 -3069.2554
## [8,]  763.6156 1373.6485  725.3805 1555.9220 1399.0801  1269.7613
## [9,] -4070.5740 1262.3947 -1880.5719 1341.1359 1408.6822  1270.2287
## [10,]  762.1407 1279.3938  806.7408 1726.5715 1479.4782  1355.4573
##           [,7]      [,8]      [,9]      [,10]
## [1,]  554.9157  763.6156 -4070.5740  762.1407
## [2,] -2286.1969 1373.6485  1262.3947 1279.3938
## [3,]  554.8381  725.3805 -1880.5719  806.7408
## [4,]  738.0696 1555.9220  1341.1359 1726.5715
## [5,]  711.8100 1399.0801  1408.6822 1479.4782
## [6,] -3069.2554 1269.7613  1270.2287 1355.4573
## [7,]  511.5332  727.1693 -4948.3299  750.4688
## [8,]  727.1693  533.4899  1316.5059 1880.2777
## [9,] -4948.3299 1316.5059  623.4179 1210.7870
## [10,]  750.4688 1880.2777  1210.7870  533.6553
```

Now let us compare the eigenvalues of each

```
eigen.XT_X = eigen(Xt_X)
eigen.ind_cov = eigen(ind_cov)
eigen.ind_cor = eigen(cor(X_std))
eigen.ind_cor_x = eigen(ind_cor_x)
svd.X = svd(X_std)

color = function(x) {
  a_func = function(a) {
    if(a == 0) {
      return("black")
    } else if(a == 1) {
      return("red")
    } else {
      return("blue")
    }
  }
  sapply(x,FUN=a_func)
}

(svd.X$d^2 / (n - 1))[1:10]


## [1] 6792.0129 1797.7053  498.3740  495.2265  493.4469  491.1551  490.8398
## [8]  490.4760  488.6024  487.0333

eigen.XT_X$values[1:10]


## [1] 6792.0129 1797.7053  498.3740  495.2265  493.4469  491.1551  490.8398
## [8]  490.4760  488.6024  487.0333
```

```
plot(svd.X$d^2 / (n - 1),eigen.XT_X$values,main="singular values^2 / (n-1) = eigenvalues")
```


singular values² / (n-1) = eigenvalues



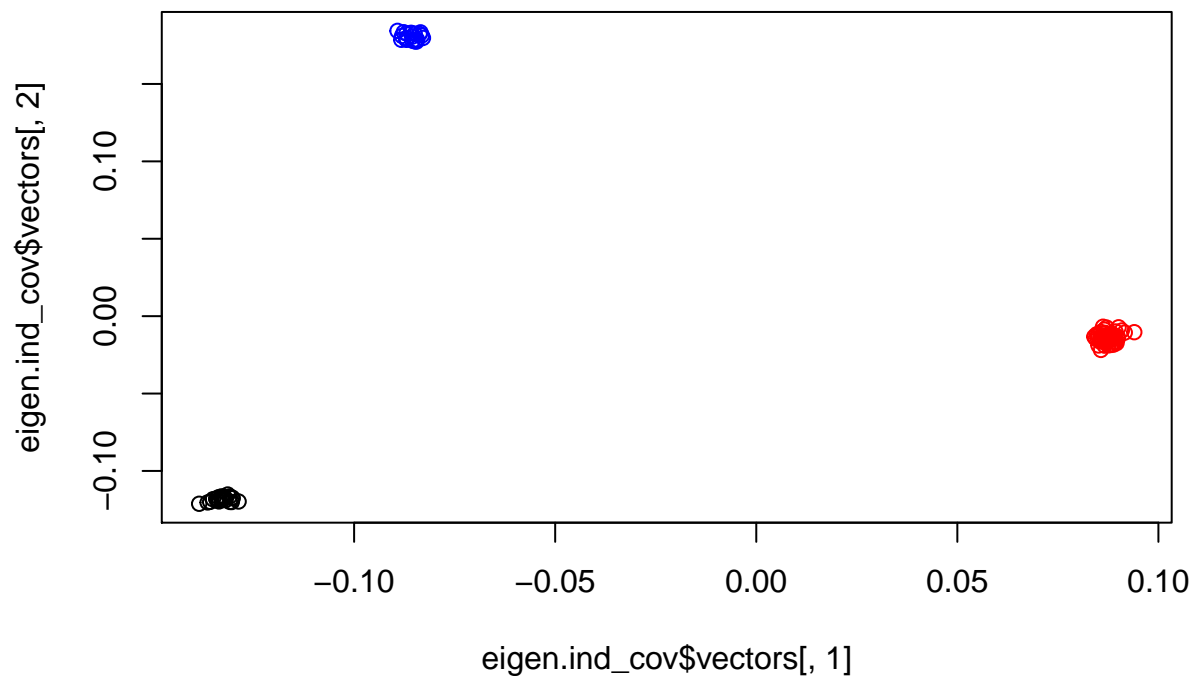
```
plot(svd.X$v[,1],svd.X$v[,2],col=color(z))
```




```
plot(eigen.XT_X$vectors[,1],eigen.XT_X$vectors[,2],col=color(z))
```




```
plot(eigen.ind_cov$vector[,1],eigen.ind_cov$vector[,2],col=color(z))
```




```
plot(eigen.ind_cor$vector[,1],eigen.ind_cor$vector[,2],col=color(z))
```




```
plot(eigen.ind_cor_x$variables[,1],eigen.ind_cor_x$variables[,2],col=color(z))
```



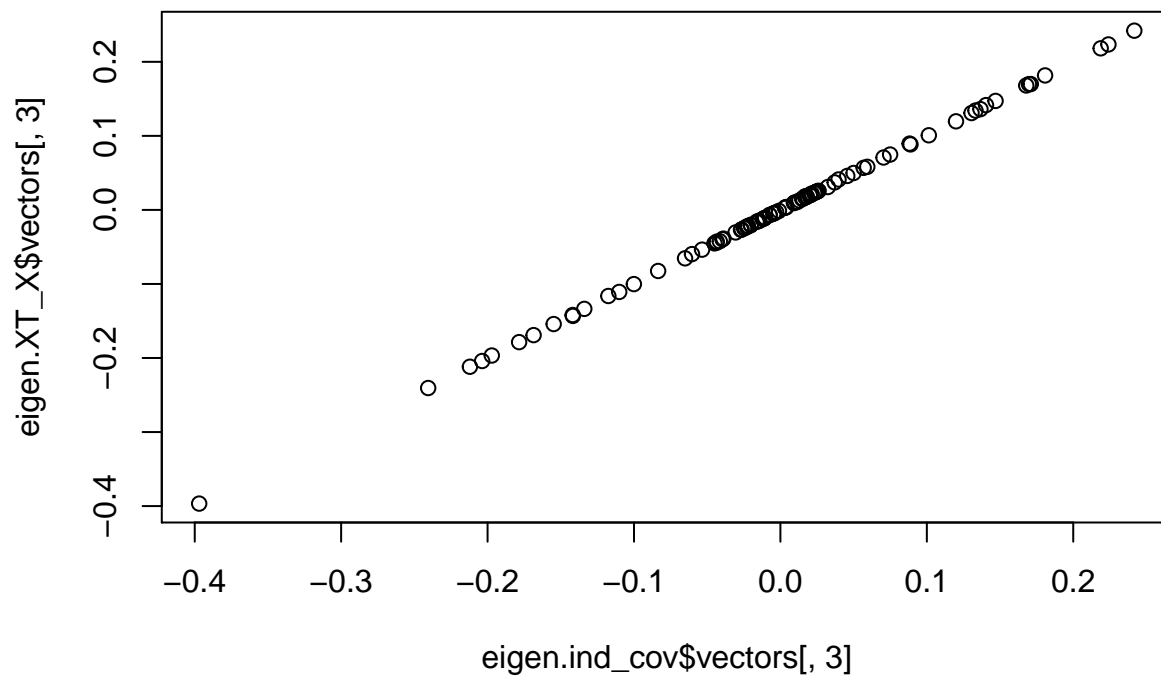
```
plot(eigen.ind_cov$variables[,1],eigen.XT_X$variables[,1])
```


```
plot(eigen.ind_cov$vector[,2],eigen.XT_X$vector[,2])
```



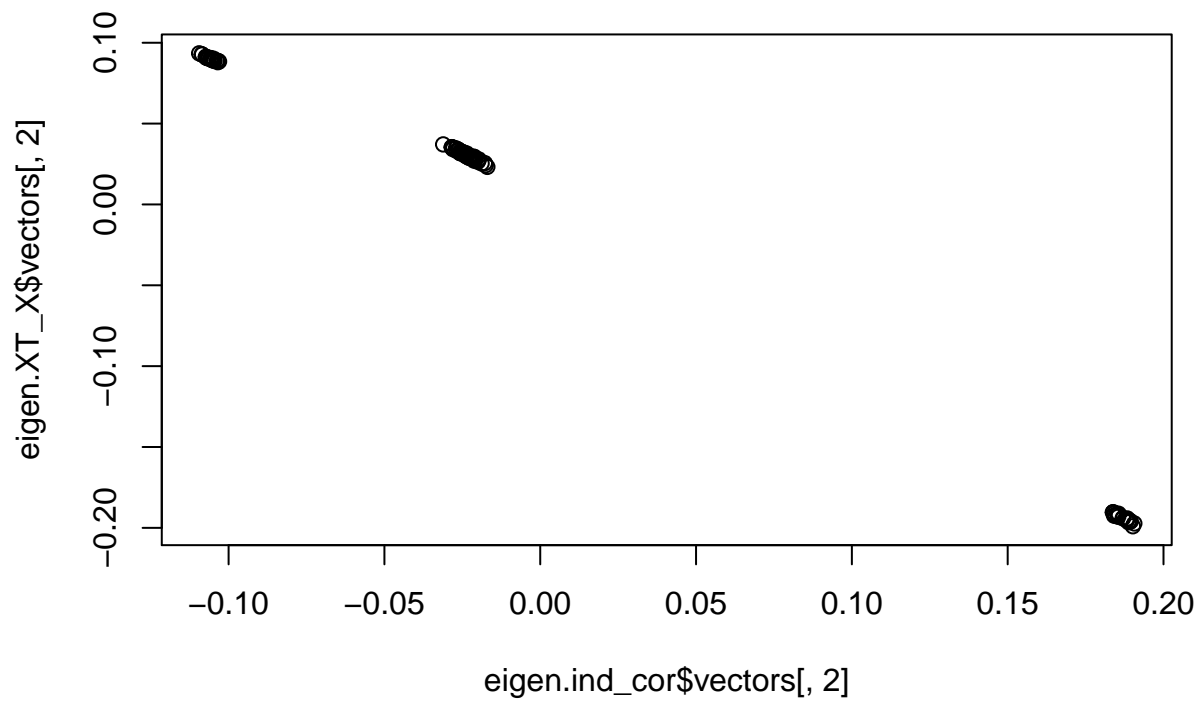
```
plot(eigen.ind_cov$vector[,3],eigen.XT_X$vector[,3])
```



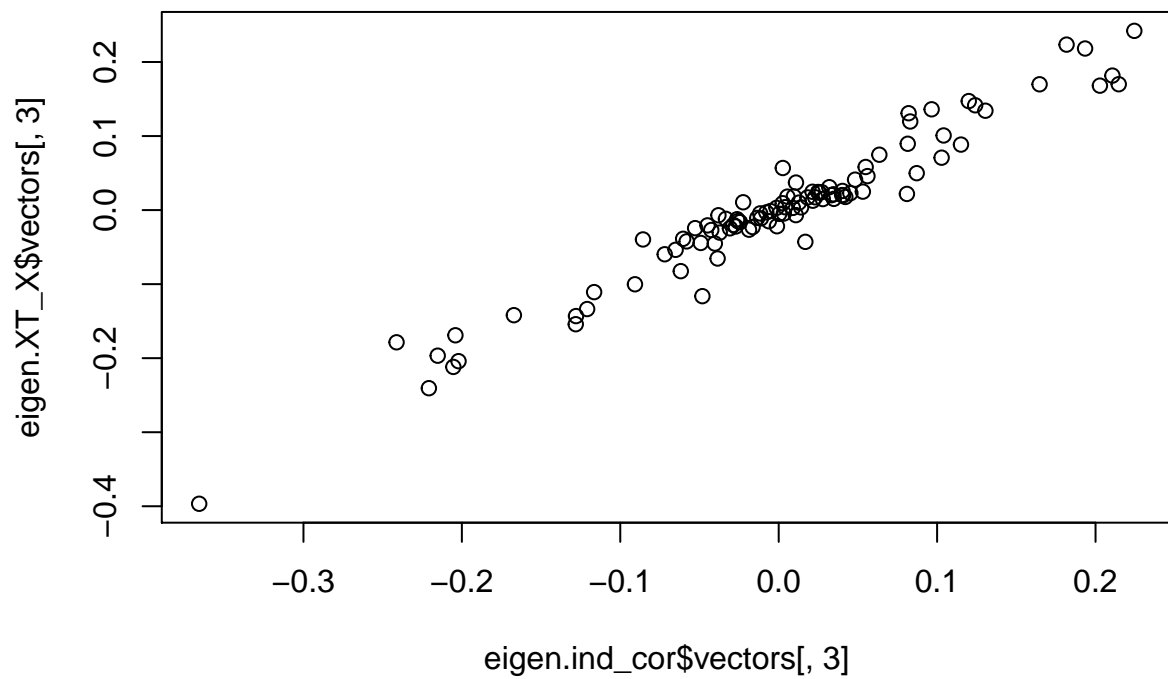
```
plot(eigen.ind_cov$vector[,1],eigen.XT_X$vector[,1])
```




```
plot(eigen.ind_cov$vector[,2],eigen.XT_X$vector[,2])
```




```
plot(eigen.ind_cor$vector[,3],eigen.XT_X$vector[,3])
```




```
plot(eigen.ind_cov$values,eigen.XT_X$values)
```



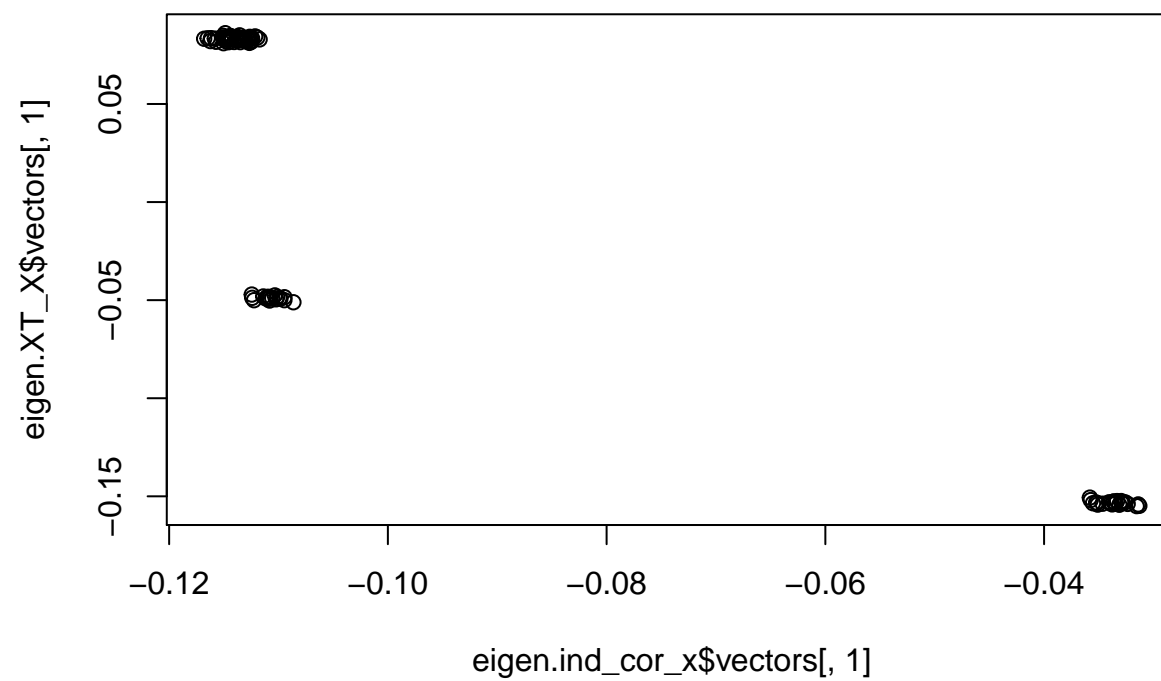
```
plot(eigen.ind_cov$values,eigen.ind_cor$values)
```




```
plot(eigen.ind_cor$values,eigen.ind_cor_x$values)
```



```
plot(eigen.ind_cor_x$vectors[,1],eigen.XT_X$vectors[,1])
```



```
plot(eigen.ind_cor_x$vectors[,2],eigen.XT_X$vectors[,2])
```



eigenvalues from svd match those from the eigenvalues from the matrix $X^T X$

Now let us see how

The